

Aman Priyanshu

Post-Training Expert | Foundation Models for Security

amanpriyanshu.github.io @ amanpriyanshusms2001@gmail.com github.com/AmanPriyanshu
Google Scholar twitter.com/AmanPriyanshu6 linkedin.com/in/Aman-Priyanshu

Professional Experience

Present Jan 2025	Foundation-AI (Cisco Systems, Inc.) AI Researcher Developed LLMs for security focused on reasoning, recursive tool usage, and agentic behavior. Specialized in post-training, test-time scaling, & bandit-style diverse path exploration for automated pen-testing.	In-Person / San Francisco, CA, USA
Aug 2024 Jun 2024	Robust Intelligence AI Security Research Intern Jailbroke LLaMA-3.1(499x) & OpenAI(4.25x Attack Success Rate) in 24h [received media coverage]; Developed automated prompt-injections; Created million-scale harmful intent dataset for AI Safety.	In-Person / San Francisco, CA, USA
Mar 2024 Jan 2024	OpenAI Red Teaming Network, Independent Contractor Participated in OpenAI led red teaming efforts to assess the risks and safety profile of OpenAI models. (NDA)	Remote / San Francisco, CA, USA
Oct 2024 Dec 2025	Anthropic Bug Bounty Program, Independent Contractor Participated in Anthropic's Bug Bounty Program (NDA). Discontinued following graduation.	Remote / San Francisco, CA, USA

Education

Dec 2024 Aug 2023	Carnegie Mellon University MSIT — Privacy Engineering	Pittsburgh, PA, USA
May 2023 Jul 2019	Manipal Institute Of Technology, MAHE B.Tech Information Technology with Minors in Big Data Analytics	Karnataka, India

Research Experience

May 2024 Aug 2023	Privacy Engineering Research [🔗] Independent Study / Advisor: <i>Professor Norman Sadeh</i> Project: For prompt-engineering geared towards usable privacy & security.	Carnegie Mellon University, USA
Aug 2023 Mar 2023	OpenMined Research Team [🔗] Project Lead and Collaborator / Collaborators: <i>Dr. Niloofar Miresghallah</i> Project: The impact of epsilon differential privacy on LLM hallucinations.	Remote / United Kingdom

Honours and Awards

- Spark Grant Winner, NOVA Hackathon, Mar 2024
- Theme Category Winner, HackCMU, Sept 2023
- Second Runners-Up - ShowYourSkill (Coursera), Jun 2022
- AAAI Undergraduate Consortium Scholar, Feb 2023
- First Prize - HackRx by Bajaj Finserv, July 2021
- First Prize - ACM UCM Datathon, UC Merced, May 2021

Publications

S=In Submission, J=Journal, C=Conference, (* = Equal Contribution)

[TR.1]	Antares: Foundation Models for Agentic Vulnerability Localization [Preprint] <i>Built Antares, a family of 350M-3B models matching GPT-5.5 on vulnerability localization at 170x lower cost.</i>	[arXiv'26]
[TR.2]	Llama-3.1-FoundationAI-SecurityLLM-Reasoning-8B Technical Report [Preprint] <i>Foundation-Sec Reasoning: first open-source cybersecurity reasoning model, rivaling 70B models at 8B.</i>	[arXiv'26]
[C.1]	When Neutral Summaries are not that Neutral: Quantifying Political Neutrality in LLM-Generated Summaries <i>2025 AAAI Conference on Artificial Intelligence</i>	[AAAI'25]
[S.1]	Are Chatbots Ready for Privacy-Sensitive Applications? An Investigation into Input Regurgitation and Prompt-Induced Sanitization [Preprint] <i>[In Submission]</i>	

Other venues of acceptances: AI4SG@AAAI'23, UpML@ICML'22, IEEE S&P'21, RCV@CVPR'21, and W-NUT@EMNLP'21.

Skills

Programming Languages	Python, Java, Go, C++, C, C#, SQL, Shell Scripting (Git & Bash)
Frameworks & Libraries	PyTorch, Tensorflow, JAX, transformers, vLLM, AdaptKeyBERT & NERDA-Con (self-authored)
Relevant Coursework	Prompt Engineering (17730), AI Governance (17716), Deep Learning (11785), Computer Technology Law (17562), Differential Privacy (17731), Information Security (17631), & Usability (17734).