

Aman Priyanshu

Post-Training Expert | Foundation Models for Security

🌐 amanpriyanshu.github.io @ amanpriyanshusms2001@gmail.com 🌐 github.com/AmanPriyanshu
🎓 Google Scholar 🐦 twitter.com/AmanPriyanshu6 in linkedin.com/in/Aman-Priyanshu

Professional Experience

Present	Foundation-AI (Cisco Systems, Inc.) AI Researcher	In-Person / San Francisco, CA, USA
Jan 2025	Developed LLMs for security focused on reasoning, recursive tool usage, and agentic behavior. Specialized in post-training, test-time scaling, & bandit-style diverse path exploration for automated pen-testing.	
Aug 2024	Robust Intelligence AI Security Research Intern	In-Person / San Francisco, CA, USA
Jun 2024	Jailbroke LLaMA-3.1(499×) & OpenAI(4.25× Attack Success Rate) in 24h [received media coverage]; Developed automated prompt-injections; Created million-scale harmful intent dataset for AI Safety.	
Mar 2024	OpenAI Red Teaming Network, Independent Contractor	Remote / San Francisco, CA, USA
Jan 2024	Participated in OpenAI led red teaming efforts to assess the risks and safety profile of OpenAI models. (NDA)	
Oct 2024	Anthropic Bug Bounty Program, Independent Contractor	Remote / San Francisco, CA, USA
Dec 2025	Participated in Anthropic's Bug Bounty Program (NDA). Discontinued following graduation.	

Education

Dec 2024	Carnegie Mellon University	Pittsburgh, PA, USA
Aug 2023	MSIT — Privacy Engineering	
May 2023	Manipal Institute Of Technology, MAHE	Karnataka, India
Jul 2019	B.Tech Information Technology with Minors in Big Data Analytics	

Research Experience

May 2024	Privacy Engineering Research [🌐]	Carnegie Mellon University, USA
Aug 2023	Independent Study Advisor: <i>Professor Norman Sadeh</i> Project: For prompt-engineering geared towards usable privacy & security.	
Aug 2023	OpenMined Research Team [🌐]	Remote / United Kingdom
Mar 2023	Project Lead and Collaborator Collaborators: <i>Dr. Niloofar Mireshghallah</i> Project: The impact of epsilon differential privacy on LLM hallucinations.	

Honours and Awards

- > Spark Grant Winner, NOVA Hackathon, Mar 2024
- > Theme Category Winner, HackCMU, Sept 2023
- > Second Runners-Up - ShowYourSkill (Coursera), Jun 2022
- > AAAI Undergraduate Consortium Scholar, Feb 2023
- > First Prize - HackRx by Bajaj Finserv, July 2021
- > First Prize - ACM UCM Datathon, UC Merced, May 2021

Publications

S=In Submission, J=Journal, C=Conference, (* = Equal Contribution)

- [TR.1] **Antares: Foundation Models for Agentic Vulnerability Localization** [Preprint]
Built Antares, a family of 350M-3B models matching GPT-5.5 on vulnerability localization at 170× lower cost. [arXiv'26]
- [TR.2] **Llama-3.1-FoundationAI-SecurityLLM-Reasoning-8B Technical Report** [Preprint]
Foundation-Sec Reasoning: first open-source cybersecurity reasoning model, rivaling 70B models at 8B. [arXiv'26]
- [C.1] **When Neutral Summaries are not that Neutral: Quantifying Political Neutrality in LLM-Generated Summaries**
2025 AAAI Conference on Artificial Intelligence [AAAI'25]
- [S.1] **Et Tu, Brute? Economic Misalignment in Personal AI Agents** [Preprint]
[Under Review | Received media coverage in Bloomberg and Inc.]

Other venues of acceptances: AI4SG@AAAI'23, UpML@ICML'22, IEEE S&P'21, RCV@CVPR'21, and W-NUT@EMNLP'21.

Skills

Programming Languages	Python, Java, Go, C++, C, C#, SQL, Shell Scripting (Git & Bash)
Frameworks & Libraries	PyTorch, Tensorflow, JAX, transformers, vLLM, AdaptKeyBERT & NERDA-Con (self-authored)
Relevant Coursework	Prompt Engineering (17730), AI Governance (17716), Deep Learning (11785), Computer Technology Law (17562), Differential Privacy (17731), Information Security (17631), & Usability (17734).